



# Database Systems

---

## Query Processing



# Last Lecture...

---

- Indexes
- Any questions?



# This Lecture...

---

- Query Processing: Overview
- What?
- Why?
- How? (Basics)



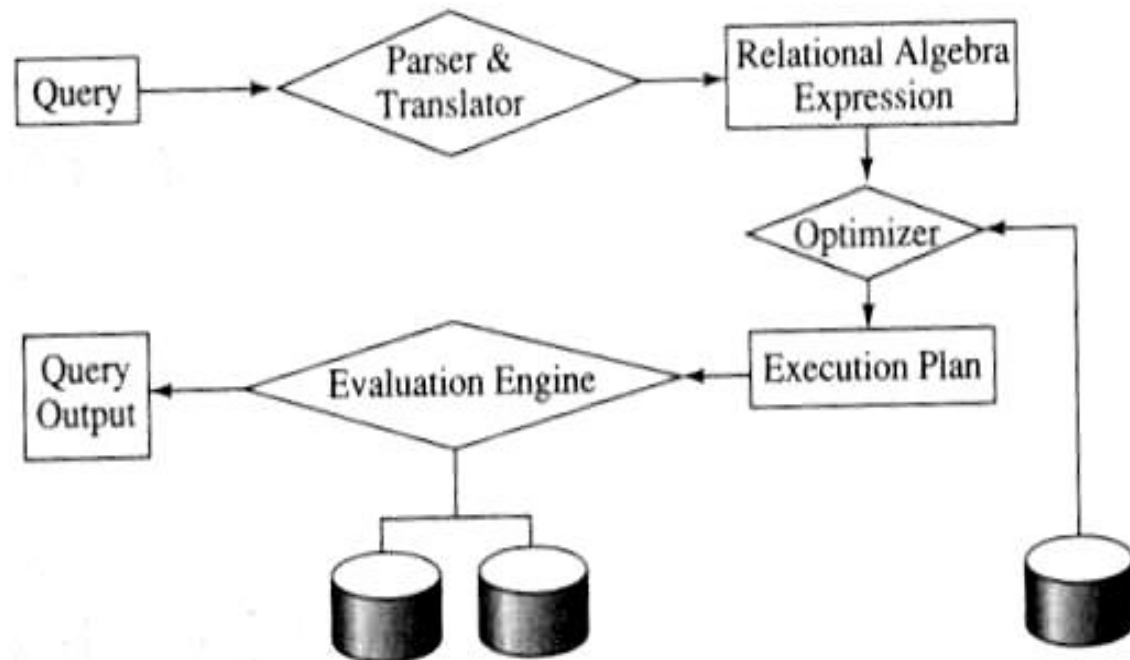
# Query Processing

---

- What happens to a query inside the DBMS?
- Query processing considers this issue.

# Steps in QP...

- The steps involved in query processing include...
  - Parsing and translation
  - Optimization
  - Evaluation





## Steps in QP... (contd.)

---

- Parsing & translation step verifies the correctness of the query and converts the query into an internal form (usually an extended relational algebra expression)
- This internal form is either a tree (i.e. query tree) or a graph (i.e. query graph)



## Steps in QP... (contd.)

---

- Next, an efficient execution strategy for retrieving the results of the query from the database files is generated. This step is called query optimization.
- This is the heart of query processing in a relational database system
- The evaluation engine executes the query according to the chosen plan



# Parsing & Translating...

---

- The first step in query processing is to convert the query into an form that can be executed
- Consider the following schema
  - S(sno, sname, status, city)
  - P(pno, pname, colour, weight)
  - SP(sno, pno, qty)






# Parsing & Translating... (contd.)

---

```
SELECT    s.sname
FROM      S, SP
WHERE     S.sno = SP.sno AND
          SP.pno = 'P2'
```

We can express this query as a relational algebra as follows...

# Parsing & Translating... (contd.)





# Parsing & Translating... (contd.)

---

- Usually, a SELECT-FROM-WHERE-GROUP BY called a ***query block*** is converted an extended relational algebra expression
- There can be many query blocks (i.e. with nested queries) in a complex query



# Why do we need Query Optimization?

---


- Consider the following number of tuples for S and SP

S	–	100 pages
SP	-	10000 pages
- Cost for computing Cartesian Product:
  - read  $10,000 * 100$  pages
  - write 1,000,000 pages (intermediate result)
- Selection
  - read 1,000,000 pages
  - keep 50 tuples (assuming 50 tuple SP.pno = 'P2')
- Total number of disk I/O  $\sim 3,000,000$  disk I/Os

# Optimization? (contd.)

- Consider the following relational algebra, which produces the same result
- Selection operator
  - Read 10,000 pages
  - Keep the result, 50 tuples in memory
- Join with S
  - Read 100 pages

Total cost 10,100 disk I/Os





# Heuristic Optimization

---

- A query can have many equivalent query trees
- Optimizer must find an efficient query plan to execute
- Heuristic rules are used for algebraic optimization



# Equivalence Rules...

---

- To transform a relational algebra expression from to an equivalent efficient query expression, certain equivalence rules are used.



## Equivalence Rules... (contd.)

---

- Conjunctive selection operations can be deconstructed into a sequence of individual selections. This transformation is referred to as a cascade of  $\sigma$ .

$$\sigma_{C1 \wedge C2}(E) = \sigma_{C1} (\sigma_{C2} (E))$$

- Selection operations are commutative.

$$\sigma_{C1} (\sigma_{C2} (E)) = \sigma_{C2} (\sigma_{C1} (E))$$

- Cascade of  $\Pi$ .

$$\Pi_{L1} (\Pi_{L2} (\dots (\Pi_{Ln} (E) \dots))) = \Pi_{L1} (E)$$





# Equivalence Rules... (contd.)

---

- Selections can be combined with Cartesian products and theta joins.

a.  $\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$

- This expression is just the definition of the theta join.

b.  $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

Etc.



# Equivalence Rules (contd.)

---

- Equivalence rules states that two expressions are equal
- It does not state, which one is better
- A large number of plans are possible! Estimating the cost for each is prohibitively expensive
- The optimizer uses heuristic rules to prune the plan space (reduce the number of plans to be considered)
  - E.g. Bring selects down, avoid cartesian products, etc.



# Indexes and cost of query plans...

---

- Using an index does not necessarily mean efficient query plan.
- Can you think of an instance where using an index is inefficient?
- Query optimizer estimates costs to compare different execution plans



# Cost estimation

---

- Execution plan has
  - A set of relational algebra operators (query plan) to obtain the result of the query
  - Algorithm used to evaluate each relational algebra operators
- Some relational algebra operators have many possible ways (algorithms).
- Costs may differ significantly based on the chosen algorithm
- We will study algorithms some algorithms used for simple selections and joins



# Schema for Examples

---

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

## Reserves:

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.

## ■ Sailors:

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.



# Simple Selections

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname < 'C%'
```

- Of the form  $\sigma_{R.attr \text{ op } value} (R)$
- Size of result approximated as *size of R \* reduction factor*; we will consider how to estimate reduction factors later.
- **With no index, unsorted:** Must essentially scan the whole relation; **cost is M** (#pages in R).
- With an index on selection attribute: Use index to find qualifying data entries, then retrieve corresponding data records. (Hash index useful only for equality selections.)

# Using an Index for Selections



---

- Cost depends on #qualifying tuples, and clustering.
  - Cost of finding qualifying data entries (typically small) plus cost of retrieving records (could be large w/o clustering).
  - In example, assuming uniform distribution of names, about 10% of tuples qualify (100 pages, 10000 tuples)  
With a clustered index, cost is little more than 100 I/Os; if unclustered, upto 10000 I/Os!

# Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- In algebra:  $R \bowtie S$ . Common! Must be carefully optimized.  $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient.
- Assume:  $M$  tuples in  $R$ ,  $p_R$  tuples per page,  $N$  tuples in  $S$ ,  $p_S$  tuples per page.
  - In our examples,  $R$  is Reserves and  $S$  is Sailors.
- *Cost metric*: # of I/Os. We will ignore output costs.





# Simple Nested Loops Join

---

```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```


- For each tuple in the *outer* relation R, we scan the entire *inner* relation S.
  - Cost:  $M + p_R * M * N = 1000 + 100 * 1000 * 500$  I/Os.

# Page-oriented Nested Loop Join

- Page-oriented Nested Loops join: For each *page* of R, get each *page* of S, and write out matching pairs of tuples  $\langle r, s \rangle$ , where r is in R-page and S is in S-page.
  - Cost:  $M + M*N = 1000 + 1000*500$
  - If smaller relation (S) is outer, cost =  $500 + 500*1000$

# Block Nested Loops Join

- Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold ``block'' of outer R.
- For each matching tuple  $r$  in R-block,  $s$  in S-page, add  $\langle r, s \rangle$  to result. Then read next R-block, scan S, etc.



# Examples of Block Nested Loops

- Cost: Scan of outer + #outer blocks \* scan of inner
  - #outer blocks =  $\lceil \# \text{ of pages of outer} / \text{blocksize} \rceil$
  - Block size = available buffers - 2
- With Reserves (R) as outer, and 100 pages of R:
  - Cost of scanning R is 1000 I/Os; a total of 10 *blocks*.
  - Per block of R, we scan Sailors (S); 10\*500 I/Os.
  - If space for just 90 pages of R, we would scan S 12 times.
- With 100-page block of Sailors as outer:
  - Cost of scanning S is 500 I/Os; a total of 5 blocks.
  - Per block of S, we scan Reserves; 5\*1000 I/Os.



# Index Nested Loops Join

```
foreach tuple r in R do
```

```
    foreach tuple s in S where  $r_i == s_j$  do
```

```
        add  $\langle r, s \rangle$  to result
```


- If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.
  - Cost:  $M + (M * p_R) * \text{cost of finding matching S tuples}$
- For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding S tuples (assuming Alt. (2) or (3) for data entries) depends on clustering.
  - Clustered index: 1 I/O (typical), unclustered: upto 1 I/O per matching S tuple.



# Sort-Merge Join $(R \bowtie_{i=j} S)$

- Sort R and S on the join column, then scan them to do a “merge” (on join col.), and output result tuples.
  - Advance scan of R until current R-tuple  $\geq$  current S tuple, then advance scan of S until current S-tuple  $\geq$  current R tuple; do this until current R tuple = current S tuple.
  - At this point, all R tuples with same value in  $R_i$  (*current R group*) and all S tuples with same value in  $S_j$  (*current S group*) match; output  $\langle r, s \rangle$  for all pairs of such tuples.
  - Then resume scanning R and S.

# Example of Sort-Merge Join



<u>sid</u>	sname	rating	age	<u>sid</u>	<u>bid</u>	<u>day</u>	rname
22	dustin	7	45.0	28	103	12/4/96	guppy
28	yuppy	9	35.0	28	103	11/3/96	yuppy
31	lubber	8	55.5	31	101	10/10/96	dustin
44	guppy	5	35.0	31	102	10/12/96	lubber
58	rusty	10	35.0	31	101	10/11/96	lubber
				58	103	11/12/96	dustin

- R is scanned once; each S group is scanned once per matching R tuple. (Multiple scans of an S group are likely to find needed pages in buffer.)
- Cost:  $O(M \log M) + O(N \log N) + (M+N)$ 
  - The cost of scanning,  $M+N$ , could be  $M*N$  (very unlikely!)



# Cost-based Optimization

---

- The fact that there are many algorithms for relational operators, choosing a good query plan using heuristics is not enough
- We can calculate the cost of each candidate query tree with the possible algorithms for each operator (& the difference can be significant)
- To compare such query plans, cost-based optimization techniques are used





# Cost Estimation

---

- For each plan considered, must estimate cost:
  - Must *estimate cost* of each operation in plan tree.
    - Depends on input cardinalities.
    - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
  - Must *estimate size of result* for each operation in tree!
    - Use information about the input relations.
    - For selections and joins, assume independence of predicates.



# Statistics and Catalogs

---

- Need information about the relations and indexes involved. *Catalogs* typically contain at least:
  - # tuples (NTuples) and # pages (NPages) for each relation.
  - # distinct key values (NKeys) and NPages for each index.
  - Index height, low/high key values (Low/High) for each tree index.



# Statistics and Catalogs

---

- Catalogs updated periodically.
  - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- More detailed information (e.g., histograms of the values in some field) are sometimes stored.

# Size Estimation and Reduction

## Factors

```
SELECT attribute list
FROM relation list
WHERE term1 AND ... AND termk
```

- Consider a query block:
- Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause.
- *Reduction factor (RF)* associated with each *term* reflects the impact of the *term* in reducing result size. *Result cardinality* = Max # tuples \* product of all RF's.
  - Implicit assumption that *terms* are independent!

### System R:

- Term *col=value* has RF  $1/NKeys(I)$ , given index I on *col*
- Term *col1=col2* has RF  $1/MAX(NKeys(I1), NKeys(I2))$
- Term *col>value* has RF  $(High(I)-value)/(High(I)-Low(I))$



# Summary

---

- Query Processing consists of several steps
- Query optimization is an important task in a relational DBMS.
- Must understand optimization in order to understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).
- Two parts to optimizing a query:
  - Consider a set of alternative plans (*heuristics are used to reduce the number of possible plans to consider*).
  - Must estimate cost of each plan that is considered.
    - Must estimate size of result and cost for each plan node.
    - *Key issues*: Statistics, indexes, operator implementations.